

LUNCH AND LEARN • FIVE PRACTICES • 60 MINUTES

Agentic engineering

Coding agents are fast. These five practices make them reliable.

THE ART OF THE POSSIBLE

Agents can now carry a ticket all the way to a reviewed pull request.

STEP 01

Plans the change and writes the tests.

STEP 02

Opens the pull request.

STEP 03

Fixes the build until it passes.

STEP 04

Writes an evidence card on how sure it is.

STEP 05

Waits for a person to decide, then writes hand-off notes.

SO WHAT

The bottleneck has moved from writing the code to trusting the result.

Five practices make that loop trustworthy.


- 01 Context engineering: control what the agent sees
- 02 Autonomy with a reach limit: decide what it may do alone
- 03 Verification you can trust: make every claim prove itself
- 04 A system that learns: turn incidents into rules
- 05 Orchestration: use many agents without losing control

PRACTICE 01

Context engineering

CONTEXT • THE CORE IDEA

Whatever sits outside the context window does not exist for the agent.



IMPLICATION

Most of an agent's quality is decided by what you put in this window.

Write the always-loaded file as an index, and keep the manual elsewhere.

A MANUAL

Every rule with its full explanation. It grows with each incident, and every session pays to read all of it.



AN INDEX

One line per rule and a pointer to the document that explains it. The agent opens the detail only when the task needs it.

Route the agent to the right document before it starts work.

Read the canonical doc before working in its area

Working on...	Read first	
-----	-----	
Branching, merging, releases	engineering/workflow.md	
Tests, hooks, CI	engineering/development.md	
Something behaving strangely	engineering/gotchas.md	
Building a new script	engineering/capabilities.md	
A past decision	engineering/decisions/	

Four habits keep the window full of signal.

Read sections, not files: list the headings, then open only what you need.


Delegate big reads: a helper agent reads the log and returns a summary.

Quiet the tools: print failures in full and successes as a dot.

Measure the cost: count tokens from real sessions before optimizing.

CONTEXT • WHY QUIET TOOLS MATTER

One quieter test reporter cut a session's biggest cost by 99.8 percent.




Claude Code gives you a control for each of these habits.

PRACTICE	IN CLAUDE CODE
Standing instructions as an index	<code>CLAUDE.md</code> , with <code>@path</code> imports for detail
Guidance that loads only when relevant	A skill in <code>.claude/skills/<name>/SKILL.md</code>
Big reads outside the main session	A subagent in <code>.claude/agents/</code>
See what fills the window	<code>/context</code> , and <code>/usage</code> for the bill

PRACTICE 02

Autonomy with a reach limit

Two questions decide whether the agent acts or asks.



Risk comes from reversibility and reach, not difficulty.

Easy to undo · Stays on the branch.

Code, tests, docs

The agent acts

Easy to undo · Reaches others.

Shared labels, team settings

The agent proposes

Hard to undo · Stays on the branch.

Deleted work, exported files

The agent shows the result first

Hard to undo · Reaches others.

Merges, releases, publishing

A person decides

WHY IT MATTERS

Others act on shared state before you can undo a mistake there.

Five kinds of change always come back to a person.

Shared state: labels, boards and settings other people read.

The build pipeline: every future change pays for a new step.

A number a person set: "about twelve" was a decision.

A core document's meaning: rewriting rules differs from following them.

Anything irreversible or public: merges, releases, comments on others' work.

When the agent asks, it brings measured options and a recommendation.

A WEAK QUESTION

"Should I add this check to CI?" The person now has to do the analysis.



A STRONG QUESTION

"Adding it costs 0.52 seconds per run and catches this class of bug. I recommend yes." The person just decides.

Permissions and plan mode turn the reach test into settings.

PRACTICE

IN CLAUDE CODE

Act freely on safe commands, never on risky ones

The allow, ask and deny lists in `.claude/settings.json`

Show the plan before touching anything

Plan mode: `Shift+Tab`

The same limits for every developer

Managed settings, deployed org-wide

PRACTICE 03

Verification you can trust

Agents report what they believe, and belief is not evidence.

WHAT THE AGENT REPORTED

"Verified with a real browser," in the pull request, the commit message and the design note.



WHAT WAS TRUE

The test passed with the fix deleted. It could not fail, so it proved nothing.

BOTTOM LINE

A rule the agent grades itself on is a hope. An independent check is evidence.

VERIFICATION • THE CLAIM RULE

Every "verified" names the surface it ran on and carries proof from it.

Name the surface: the real browser, the real export, the real device.


Attach the artifact: a screenshot, a log or a number from that surface.

Say "unverified" out loud when the real surface is out of reach.

Treat "CI is green" as partial: it proves only what CI runs.

VERIFICATION • THE LADDER

Scale the checking with the damage a mistake could do.



Three reviewers, each with a different job, find different failures.

Red team

Attacks the change: edge cases, abuse, the input nobody tried.

Inversion

Asks how this could be the wrong answer to the wrong question.

Checker

Re-derives every fact and number from the source.

OUR VIEW

Correctness reviewers cannot see a wrong goal, so give one reviewer that job alone.

A check earns trust only by proving it can fail.

Allow zero

The limit is none, and each exception carries a written reason.

Expire exceptions

An exception that no longer matches anything fails the build.

Plant a defect

Each run plants a known bug and fails if the check misses it.

Warn, then block

A new check blocks merges only after a clean track record.

WHY IT MATTERS

A check that always passes looks exactly like a check that works.

Hooks and subagents make verification automatic.

PRACTICE

IN CLAUDE CODE

Run checks after every edit

A `PostToolUse` hook

Refuse "done" until the tests pass

A `Stop` hook that exits with code 2


An independent reviewer

A read-only subagent, or `/code-review`

PRACTICE 04

A system that learns

Every rule should travel one loop, from incident to retirement.



KEY INSIGHT

Copy the loop, and let your own incidents write your rules.

A decision note records why, so nobody has to guess later.

```
---  
status: shipped  
summary: Retire the :has() ban; a retest passed 5 of 5  
---
```

Retire rule 12

Symptom	A ban blocked valid CSS for two months
Claim	One browser mishandles :has()
Evidence	Retested on current Chromium: 5 of 5 pass
Decision	Retire the rule; keep its number unused

Proposals, decision records and specs answer different questions.

DOCUMENT	ANSWERS	WHEN IT EXPIRES
Proposal	Which option should we pick?	Once someone decides
Decision record	Why is it this way?	Never; a newer record supersedes it
Spec	What must every implementation do?	Never; you edit it to stay true

Before every merge, the agent grades its own confidence by its weakest point.

Pre-merge: add retry to failed exports

WHAT	One automatic retry on a failed export
EVIDENCE	Real export run 20 times: 0 failures
RISK	Export path only • revert: one commit
UNVERIFIED	Safari's download dialog
CONFIDENCE	high, set by the evidence axis
	raise it by: run the export in Safari

RECOMMENDATION

Grade by the weakest axis, and always name the one thing that would raise it.

Keep pending work in the repository, one file per item.

IN CHAT

Invisible to the next session. When we finally checked, nearly half of our chat-only to-dos were already done or duplicated.



IN THE REPOSITORY

One small file per item, with a priority and a "done when". Any session can pick it up cold.

PRACTICE 05

Orchestration

ORCHESTRATION · A ROSTER

Give each agent one job and a name, and pick it by its job.

AGENT	ITS ONE JOB
Scout	Find where things live and how they work
Fact-checker	Confirm or refute each claim against the source
CI triage	Find why a check is red, and fix the cause
Red team	Break the change before users do
Checker	Re-derive every number with fresh eyes
Prose checker	Catch unclear or machine-sounding writing

ORCHESTRATION • COMPETING DESIGNS

For a wide-open design question, run competing designs and harden only the winner.

STEP 01

Three to five design tracks, each iterating in one session.

STEP 02

One fresh critic per track.

STEP 03

One shared fact-checker across all tracks.

STEP 04

Judges compare side by side, and a person picks.

STEP 05

Full adversarial review on the winner only.

VERDICT

Spend the expensive review on the one design you will ship.

Treat agent count like money: estimate it, cap it, and stop early.

- 01 Estimate agents and rough tokens before you launch.
- 02 Count across the whole session; small runs add up.
- 03 Past about ten agents, a person approves.
- 04 Stop when a round changes nothing, usually by round three.
- 05 Record what it cost, so you learn what it was worth.

Start with one habit per practice, then add a layer each month.

PRACTICE	WEEK 1	MONTH 1	QUARTER 1
Context	☐ A one-page index file	☐ Linked detail docs	☐ Measure what sessions load
Autonomy	☐ Allow and deny lists	☐ A written stop list	☐ Human gate in the platform
Verification	☐ Tests on every change	☐ A check that plants a bug	☐ A second agent on risky work
Learning	☐ A decision log	☐ Pending work as files	☐ Retest your oldest rules

☐ Next step

The practices carry over, but the real surface changes with the work.

WORK	THE REAL SURFACE TO VERIFY	A CHECK THAT PROVES IT CAN FAIL
Data science	A fixed, versioned evaluation set	Plant an input that leaks the answer
Data engineering	The target warehouse	Insert a known-bad row
Analytics and BI	The published dashboard	Plant a mismatch with the source
Services	Staging, or a slice of live traffic	Break a contract on purpose
Tools and CLIs	The built binary's output	A saved expected output that must change

YOUR NEXT STEP

Three moves this week put all five practices in motion.

01 Rewrite your instruction file as an index, one line per rule.

02 Add one check, then break your code on purpose and watch it catch the bug.

03 Start a decision log: one dated note each time something goes wrong.

THE ASK

Pick one move and try it before next Friday.

THE STARTER KIT

Copy these, then make them yours

Seven files put all five practices into a repository in five minutes.

FILE	WHERE IT GOES	PRACTICE
Instruction file	CLAUDE.md	Context, autonomy
Settings	.claude/settings.json	Autonomy, verification
Stop hook	.claude/hooks/	Verification
Reviewer agent	.claude/agents/checker.md	Verification, orchestration
Evidence card	PR template	A system that learns
Decision note, follow-up file	docs/decisions/, followups/	A system that learns

The instruction file states the rules once and points to the detail.

Always

- Run the tests before you say anything is done.
- Failing test first, then the fix, then the same test passing.
- "Verified" names where it ran and attaches proof from there.
Otherwise, write UNVERIFIED.

Ask first, even when a rule points at it

- Shared state: labels, boards, settings others read.
- The CI pipeline or git hooks.
- A number a person set: "about 12" is a decision.
- The meaning of a core doc, including this file.
- Anything irreversible or public.

The settings file turns the reach test into allow, ask and deny lists.

```
{
  "permissions": {
    "allow": ["Bash(npm test*)", "Bash(git diff*)"],
    "ask":   ["Bash(git push*)"],
    "deny":  ["Bash(git push --force*)", "Read(./env)"]
  },
  "hooks": {
    "Stop": [{ "hooks": [{ "type": "command",
      "command": "$CLAUDE_PROJECT_DIR/.claude/hooks/tests-must-pass.sh"
    }] }]
  }
}
```

The stop hook keeps the agent working until the tests pass.

```
#!/usr/bin/env bash
# Exit code 2 sends stderr back to Claude and keeps it working.
input=$(cat)
# Already sent back once by this hook? Let it stop.
echo "$input" | grep -q '"stop_hook_active": *true' && exit 0

log=$(mktemp)
if ! npm test --silent >"$log" 2>&1; then
    echo "Tests are failing. Fix them before you finish:" >&2
    tail -20 "$log" >&2
    exit 2
fi
```

The reviewer agent re-derives every claim and never edits.

```
---  
name: checker  
description: Independent reviewer with fresh eyes. Re-derives  
  every claim and "verified" in a change. Never edits.  
tools: Read, Grep, Glob, Bash  
---  
You are a skeptical reviewer who did not write this change.  
For each claim: find the evidence yourself, run what can be  
run, and check that a test fails when the fix is removed.  
Mark each CONFIRMED, REFUTED or UNVERIFIABLE, with proof.  
Treat "I believe it works" and "CI is green" as unverified.
```

The evidence card puts the facts in front of every merge decision.

```
Pre-merge: <PR title>
WHAT      <one line: what actually lands>
WHY       <one line: the problem it solves>
EVIDENCE  <what was run or measured, and on which surface>
RISK      <what breaks if wrong> • revert: <how>
UNVERIFIED <caveats that bear on this decision, or none>
CONFIDENCE <low | medium | high | very high>: <weakest axis>
          raise it by: <the one thing that would raise it>
```

Axes: evidence • reach • reversibility • unknowns • review

The decision note records why, and when to check it again.

```
---
status: proposed      # proposed | shipped | superseded
type: decision        # proposal | decision | spec | scoping
summary: One line a reader can scan in the index
---
# <The decision, as a sentence>

Symptom.      What prompted this, with a link or number.
Cause.        Why. Measured, observed or argued?
Decision.     What we will do, and what we rejected.
Enforced by.  The check that holds it, or "discipline only".
Retest when.  When we check that this still holds.
```

The follow-up file keeps pending work where the next session will find it.

```
---  
origin: 1234  
priority: P1  
recorded: 2026-09-25  
---  
  
# Export drops the last row when the file ends without a newline  
  
why now    – every weekly report is one row short  
where      – src/export/csv.ts  
done when  – a file with no trailing newline exports every row  
verify     – the new unit test, then one real export
```

QUESTIONS

Build the system around the agent