FIX · _FOCUSSTYLE: BLUR

Blur softens the rest and moves nothing.

`blur` used to scale its target and cast a hard shadow, which shifted the cards and clipped the outer ones at the stage edge. Now it only defocuses the rest.

The corner card stays inside the stage.

Signal Intake

Weekly structured collection across customer conversations, market data, and competitive moves. Normalized into a common schema before scoring.

Scoring Model

Each signal scored on three dimensions: confidence, recency, and strategic relevance. Weights are team-configurable — configured once, then defended forever.

Decision Log

Every decision recorded with the signals that informed it, the options considered, and the criteria applied. Feeds the calibration loop.

Calibration Loop

Monthly retrospective that compares predicted outcomes to actual outcomes and adjusts scoring weights accordingly.

The lower card stays full width.

False signal amplification

One loud voice dominating the decision. The model caps any source at 30% — a ceiling one enterprise customer tests monthly.

Signal hoarding

Signals collected, decisions unlogged — the loop learns nothing. Above P2, the log is required. No log, no change.

A lower row never nudges the rows above.

- [x] Signal taxonomy ratified, in workshop four of three
- [x] Scoring weights agreed by the steering committee
- [x] Decision log live in staging
- [-] Pilot teams trained, two still "circling back"
- [-] Operating rhythm on the calendar, attendance optional in practice
- [ ] Exec sponsor confirmed for the launch comms
- [/] Per-team weighting UI, descoped to next half

The right-hand verdict keeps its edge.

Path A · The spreadsheet

Speed · Auditability · Adoption · Calibration

Open everywhere, version unknown. No trace — in fairness, the appeal.

Path B · Vendor North

Speed · Auditability · Adoption · Calibration

Clean intake, real audit trail. The weights stay on their side of the contract.

Path C · Vendor West

Speed · Auditability · Adoption · Calibration

The better demo, by some distance. The weights are theirs; so is §9.2.

Path D · Build in-house

Speed · Auditability · Adoption · Calibration

All four criteria, inside the budget, weights ours. Recommended.

Every pricing tier sits where it rests.

Vendor North

$38 / seat / mo

- Signal-intake connectors
- Calibration weights exposed
- Exit without repricing

For teams that enjoy renewals.

MOST POPULAR

Vendor West

$52 / seat / mo

- Signal-intake connectors
- Calibration weights exposed
- Exit without repricing

For teams that enjoy nicer renewals.

Build in-house

4 engineers

- Signal-intake connectors
- Calibration weights exposed
- Exit without repricing

For us, apparently, again.

Step four ends flush with the stage.

STEP 01

Pick one team and one decision type

A team with an existing rhythm, one decision category, thirty days.

STEP 02

Log everything, decide nothing differently

Month one changes nothing but the record. Decide as you always did; write it down.

STEP 03

Run your first retrospective

Day 30: score decisions against outcomes. The model's first calibration pass.

STEP 04

Expand to a second team

One retrospective in hand, onboard team two with data, not promises.

Focus without the shove.

`_focusStyle: blur` *keeps every card at its resting size and position, so* `_focusSteps` *walks a slide without anything jumping.*