




DECISION FRAMEWORK · Q3 2025

From Signal to Strategy

A 59-slide answer to the question "have you considered writing things down"



What this deck covers, in order


01 Why This Exists — slide 3

→ **The Framework — slide 7**

03 The Evaluation — slide 23


04 The Build — slide 33

05 The Results — slide 48



The window for differentiation is narrowing

Commoditized infrastructure, compressed release cycles, and rising switching costs have cut the average durable advantage from 36 months to under 14. This slide will appear in every deck for the next two years regardless of whether that stays true.




“

The signal was always there. We just didn't have a system that forced us to look at it before we'd already decided.

”

— Head of Product, Pilot Team 3, in a retrospective where we then decided what we'd already decided



IMPACT · PILOT RESULTS

Six months of results across four product teams

Same teams, same conditions, same spreadsheet.

73%

FASTER DECISION CLOSE

4.2×


SIGNAL RECALL

18

DECISIONS LOGGED

91%

TEAM ALIGNMENT



CALIBRATION RESULT • 6-MONTH PILOT

14X


Return on signal investment — on a
baseline the framework team defined.
Verified by nobody.



01

SECTION 01 • THE FRAMEWORK


We built a four-component scoring system. Two of the four are used regularly



SIGNAL DEFINITION • WORKSHOP 04


Before we score signals, we need to agree on what a signal is


Three workshops in. The definition gets
socialized in a fifth before anyone can use it.



How signals move from input to decision

Four-stage pipeline — 11 weeks in, still in pilot





The framework has four components

Signal Intake

Weekly structured collection across customer conversations, market data, and competitive moves. Everyone agrees it's a good idea. Nobody does it the week of the retrospective.

Scoring Model


Confidence, recency, and strategic relevance, each weighted. The weights are team-configurable — which means the head of product reconfigures them until the output agrees with the roadmap.

Decision Log

Every decision recorded with its signals and criteria. A required artifact. It holds 18 entries, against roughly 340 decisions actually made.

Calibration Loop

A monthly retrospective comparing predicted outcomes to real ones. The meeting reliably exists. The predictions, reliably, do not.



Signal Intake produces three outputs

1

Weekly Signal Brief

Last week's top 10 signals, ranked, with confidence scores and sources. Sent to product leads every Monday — where it sits unread in a folder called "Framework Stuff."

2


Anomaly Alerts

Real-time flags when a signal crosses the 2σ threshold, routed to the accountable PM on a 4-hour SLA. The PM usually replies inside the 4 hours — to ask what 2σ means.

3

Monthly Signal Index

The source of truth for the calibration loop, and required reading before every retrospective. Nobody has read it. It is comprehensive.



The three things the framework connects

Signal


The observed input — a verbatim, a metric move, a competitor's announcement. The unit of intake. Frequently confused with "things the VP heard at a conference," which enter the pipeline already scored.

Decision

A signal plus a deadline, logged with its rationale. Every decision is meant to trace to a signal. In practice the signal is "we discussed it at the offsite," reconstructed for the log the following week.

Outcome

The result, measured against the prediction at retrospective. The unit of calibration. 18 have been logged. Roughly 340 have occurred; the other 322 are filed under "context."



Two failure modes the framework is designed to prevent

False signal amplification.


A single loud voice — one enterprise customer, one analyst report, one VP with a feeling — dominates the decision without ever being weighed against the full signal set. The model caps any one source at 30% of total weight. Unless that source is the CEO, in which case the cap is a guideline.

Signal hoarding.

Teams collect signals but never log decisions, so the calibration loop has nothing to learn from. The rule — no log, no change above P2 — was printed on a poster and hung in the meeting room. The poster has since been replaced by a free-pizza flyer.



Scoring Model Deep Dive



What the scoring model actually does


The most configurable component — a feature or a warning sign, depending on your team. Three dimensions:

- 1

Confidence
Independent corroborating sources, 1–5. Enterprise customers count as 1, regardless of volume.
- 2


Recency
Time-decay, configurable half-life. Set to two weeks, then surprised only recent news scores.
- 3

Strategic Relevance
Owner-scored, 1–5. A 5 correlates with whoever is presenting the roadmap this quarter.



The six signal dimensions

01	Confidence	Independent sources corroborating it	1-5 • AUTO • ENTERPRISE ALWAYS GETS A 4
02	Recency	Time-decay from signal date	0.0-1.0 • AUTO • "SHORT" AFTER A BAD Q
03	Relevance	Alignment to current bets, owner-scored	1-5 • MANUAL • WHAT THE PM ALREADY PLANNED
04	Reach	Customers or segments affected	1-5 • AUTO • 5 IF THE ENTERPRISE ASKED
05	Effort	Engineering and design cost to act	1-5 • MANUAL • ENGINEERING'S EYEBROWS GO UP
06	Confidence delta	Change since the last scoring cycle	-5 TO +5 • AUTO • DID ANYONE ASK A CUSTOMER



Scoring model: before and after the calibration loop

BEFORE CALIBRATION


Equal weights — confidence, recency, and relevance each contribute 33%. Simple, and at least honest that we are basically guessing. In eighteen months no team has changed them from 33.



AFTER CALIBRATION

Weights reflect your team's historical accuracy — except the team keeps changing, and the history is three months of data from a quarter everyone agrees was atypical. So most teams keep the 33s and call them calibrated.

The shift from equal to calibrated weights takes two retrospective cycles — 60 days from adoption, or 14 months, depending on who you ask.




Two intake modes for different signal types

Structured Intake

Clear-schema signals — NPS verbatims, support tickets, win/loss notes. Ingested and scored automatically, with zero manual handling. Produces 94% of the data and 12% of the roadmap decisions — the part of the pipeline that works, which is why no one discusses it.


Unstructured Intake

No-schema signals — field notes, conference hallway talk, a board member who had thoughts at a dinner. Routed to the owner for manual classification. Produces 6% of the data and 88% of the roadmap decisions, every one of them urgent.



How a decision moves through the framework





What the framework does not do


Doesn't make decisions — it formalizes the one you'd have made anyway.

Doesn't replace discovery — it routes whatever surfaces before the roadmap locks.

Doesn't work without the Decision Log. 18 entries; ~340 decisions made.


Doesn't guarantee alignment — the same argument, just earlier in the quarter.

Doesn't scale below P2 — the daily 90% still lives in Slack threads from 2022.




Four things that must be true before you begin

- 01 A monthly prioritization cadence. "Whenever someone escalates loudly" doesn't count — it just gets louder.
- 02 Someone who owns signal collection. Which means defining a signal first. See slide 8.
- 03 Leadership logging rationale, not just outcomes. This slide exists because that part is hard.
- 04 90 minutes a week. Budget four.



How a week of signals narrows to one logged decision





SECTION 02 WEIGHS FOUR TOOLS AGAINST FOUR CRITERIA — DEFINED BY THE TEAM THAT BUILT ONE. DISCLOSED
IN A FOOTNOTE.


SECTION 01 OF 05 COMPLETE

The framework is specified —
the real question is build or buy



SECTION 02 • THE EVALUATION

We evaluated four tools.
One of them was built
by the evaluation team



Four requirements every decision system must meet

01 Speed

Decisions close inside their window. Six months to calibrate isn't fast. We won't dwell on that here.

02 Auditability


Every decision above the threshold carries a traceable rationale — for after the launch goes badly.

03 Adoption

Weekly use, or calibration never runs. We budgeted 90 minutes per PM. Actual: 11.

04 Calibration

It has to improve. A static model is a spreadsheet with a dashboard nobody checks.



We evaluated four intake tools against the criteria

Tool A • Chorus

- ✔ Speed
- ⚡ Auditability
- ✔ Adoption
- ✖ Calibration

Great call recording, no logging, no calibration. The sales team already uses it.

Tool B • Productboard

- ✖ Speed
- ✔ Auditability
- ✔ Adoption
- ✖ Calibration

Solid intake, no calibration. Setup was "3–4 weeks." It took 11.

Tool C • Notion

- ✔ Speed
- ✔ Auditability
- ⚡ Adoption
- ✖ Calibration


Build anything. Teams built seven, then debated the canon for two quarters.

Tool D • Sprig + Decision Log



- ✔ Speed
- ✔ Auditability
- ✔ Adoption
- ✔ Calibration


Built by this evaluation's authors. Meets all four — as they defined them.



The four tools side by side


CRITERION	CHORUS	PRODUCTBOARD	NOTION	SPRIG + LOG
Speed	✓	✗	✓	✓
Auditability	✗	✓	✓	✓
Adoption	✓	✓	✗	✓
Calibration	✗	✗	✗	✓
Setup time	1 day	3–4 weeks	40+ hours	Same day


✦ Criteria defined by the team building Sprig + Log. We're transparent about it. It's in the footnote.



SCALE · 0-10, ON THE CRITERIA WE WROTE

The four tools, scored across the criteria we wrote





How we sort the four tools against our two axes

Coverage · Cost

High coverage / Low cost

Sprig + Log — best coverage, lowest TCO. Built by the evaluation team.
Productboard — narrower, here to show we looked.

High coverage / High cost


Notion — full coverage in theory. We tried it: seven versions, a quarterly canon debate.

Low coverage / Low cost


Chorus — cheap, three criteria short. The sales team likes it.

Low coverage / High cost


None — either the signal or a gap. We're treating it as a signal.



The evaluation came down to one question the vendors could not answer



The left card is struck through; the DECISION connector is bold. The conclusion came first — the slide was built to hold it.



We are building, not buying

Decision · 2026 Q1

BUILD


Owens the scoring policy, the calibration loop, and the timeline — plus the maintenance, the on-call rotation, and every future explanation of why the framework scored the wrong thing.

WHY NOT BUY

Three vendors, none exposing calibration weights. The weights are the product; you cannot buy the product without them. That was the finding.

WHY NOT DELAY

The window closes in 18 months — a sentence that has been in this deck, unchanged, since Q1 2025. Each quarter we update the deck, not the window.



How we make calls when the spec is silent

- 01 Default to the choice that's cheaper to reverse — unless reversing it needs a meeting.
- 02 Name the actor, never the system. The system can't be held accountable. The actor can be reorganized.
- 03 Write the bet on the slide with the choice. Not always where it gets reviewed.

What the build covers

Three phases, four workstreams. We own the policy, the loop, the timeline — and whatever Phase 2 becomes.

1

Phase 01 — Architecture

Scoring policy live, the Decision Log accepting entries, one pilot team in weekly cadence.

2

Phase 02 — Operations

Multi-team calibration, automated weight updates, and before-after data you could defend in a board update.

3


Phase 03 — Scale

Org-wide enablement. In the roadmap since 2024. We remain committed.

T


SECTION 03 • THE BUILD

Three phases ship the architecture, the operations, and the scale — and Phase 01 has shipped





What ships in each phase

WORKSTREAM	PHASE 01	PHASE 02	PHASE 03
Signal Intake	Connector v1	Multi-source dedupe	Anomaly auto-routing
Scoring	Equal-weights model	Per-team calibration	Per-decision profiles
Decision Log	Append-only schema	Outcome auto-pairing	Examiner export
Adoption	One pilot team	—	Org-wide enablement



The build schedule — Phase 3 is the big red bar, as it was last year






WORKSTREAM BOARD · TODAY

The build board, mostly one column

BACKLOG	IN PROGRESS	SHIPPED
Examiner exportM • decision-log blocked	Per-team calibrationL • scoring at-risk Outcome auto-pairingM • decision-log Org-wide enablementXL • adoption blocked	Scoring policy v1M • scoring



Who owns each part of the framework lifecycle

Signal custody

Signal owner

Manages intake quality. Tunes weights only by choosing what to surface.

Policy

Framework operator

Owns scoring, cadence, rollback. One person. Noted in the risk register.

Consumption


Product team

Runs intake and logging. Mostly asks the operator to change the weights.

Oversight

Auditor

Reads the audit trail. Read it once, found 18 entries, asked if that was expected.



Three phases get us from decision to org-wide adoption

PHASE 01

Architecture

Scopes what we build, what we buy, and what we defer. The output is an architecture decision record — which will itself need an architecture phase before it can be approved, and a workshop to define "decision."

PHASE 02


Pilot

One team, one decision type, one quarter. The phase ends at production cadence — meaning the retrospective happened at least once, and someone screenshotted it for the board.

PHASE 03

Rollout

Five teams in two months, ending above 90% adoption. Planned for Q2 — as it has been for three consecutive years. Q2 is load-bearing.



Three milestones mark Phase 01 complete

MILESTONE A

Scoring policy in production

The signed policy runs end-to-end and the first calibrated brief lands in leadership's inbox. Their first reply asks whether the weights can be adjusted. They can — and did, that afternoon.

MILESTONE B


Per-team weights

One framework carries distinct weights per team without forks; recalibration is a single update. Each team will still want its own, and three will fork it anyway.

MILESTONE C


Per-decision-class profiles

Authoring a profile for one decision class takes minutes. Agreeing what counts as a decision class takes a workshop series. See slide 8, which has not moved.



Phase 01 acceptance — what shipped, what slipped, what stayed open

✓	Scoring policy live across all pilot teams	
✓	Decision Log audit trail readable by Auditor role	shipped 2026-Q1
✓	One reference team running weekly cadence	one team
—	Examiner pack auto-generation	was Phase 01 • now Phase 1b
○	Adoption above 90%	Phase 03
○	Culture change	not in roadmap



Decisions used to require a quarterly re-litigation

Before and after the framework

BEFORE


First principles every time. Close: 4 hours, p99 an offsite. Outcome: what the senior person wanted, called consensus.



AFTER

Resolved against logged weights. Close: 18 minutes. Outcome: what the model says — once the weights match the senior person.

The architecture change is the calibration loop. The culture change is still in Phase 02.



How to roll this out across your organization

STEP 01

Pick one team and one decision type

Start with a team that has a real rhythm. No rhythm, no framework.

STEP 02

Log everything, decide nothing differently

Just log, for a month. "Low-effort." Highest dropout rate.

STEP 03


Run your first retrospective

Day 30, score against outcomes. No outcomes? Score the room, call it "qualitative."

STEP 04

Expand to a second team

Your evidence: one team ran one retrospective. Use it confidently.



The weekly practice in three moves

STEP 01

Sense

Observed inputs, never invented. Skipped about 70% of the time.



STEP 02

Score


A signal is data once it carries a number. Calibrating it takes 14 months.



STEP 03

Decide

A signal plus a deadline. The loop closes if anyone logged a prediction.



The case for the framework in three moves

Claim


Calibrated prioritization with audit-grade decision custody. We stop paying the re-litigation cost on every quarterly review.

Evidence

Six months across four product teams. Close-time dropped; calibration ran once. The four teams have since been reorganized into three. None of them count in the "before" baseline.

Implication

The framework works. The deck says so. The deck was written by the framework team. We trust the framework team — the framework told us to.



Wiring a signal into the framework is three lines of application code


JavaScript • DecisionFramework SDK v2.847 transitive dependencies

```
import { DecisionFramework } from "@company/signal-sdk";

const framework = new DecisionFramework({ configFile: "./framework.config.json" });

// Score a signal at intake
const score = await framework.score(signal, { dimensions: ["confidence", "recency", "relevance" ]});

// Log every decision – calibration depends on it
// (nobody calls this line in production, but it is here)
const entry = await framework.decisions.log(decision, { signals: [signal.id], rationale });
```



BEFORE & AFTER • SCORING MECHANICS


Spreadsheet-driven scoring versus framework-driven scoring that is basically also a spreadsheet

BEFORE • THE HONEST SPREADSHEET

```
# Manual scoring. Auditable in the  
# sense that you can see who last  
# edited the file  
import pandas as pd  
  
signals = pd.read_csv('./signals.csv')  
signals['score'] = signals.apply(  
    lambda r: 0.33*r.confidence  
    + 0.33*r.recency + 0.33*r.relevance,  
    axis=1,  
)  
signals.to_csv('./scored.csv')
```

AFTER • THE FRAMEWORK

```
# Calibrated weights, signed policy,  
# every score is audit-logged,  
# same math as the spreadsheet  
from decision_framework import Calibrator  
  
calibrator = Calibrator.load('./policy.json')  
for signal in calibrator.intake.unscored():  
    calibrator.score(signal)  
    calibrator.decisions.log_if_relevant(signal)
```



SECTION 04 PRESENTS SIX MONTHS OF PRODUCTION DATA. THE DENOMINATOR IS CAREFULLY CHOSEN.

SECTION 03 OF 05 COMPLETE


The framework is built. Twelve
percent of eligible PMs use it



SECTION 04 • THE RESULTS


Six months of data.
Eighteen logged decisions.
One calibration cycle

The pilot team measured what the pilot team built.



Before the results, what you should have learned in the first 47 slides

- 01 Four components. Two used regularly. — see Section 01
- 02 The evaluation team recommended the tool the evaluation team built. — see Section 02
- 03 Phase 01 shipped. Phase 03 in the roadmap since 2024. — see Section 03
- 04 Adoption is 12%. Target 90%. The gap is "cultural." — see this section
- 05 The calibration loop has run once. We call it "calibrated." — see this section



Where we are against quarter targets

94%

Signal-classification success

TARGET 99%, GAP IS "KNOWN ISSUE"

18 min

p99 decision close

TARGET 20 MIN, BEATING TARGET

18


Decisions logged

TARGET 340, GAP IS "CULTURAL"

1

Calibration cycles run


TARGET 6, GAP IS "STRUCTURAL"




H1 2026 · PHASE 01 READINESS

Phase 01 readiness, by workstream

The status pills reflect the most optimistic reading of the data.




Source: Linear · "blocked" means blocked, we are working on the wording




H1 2026 · 1,840 PERSON-HOURS

Where the engineering quarter went

The "Toil and on-call" wedge is the one nobody put in the roadmap.




Last updated 2026-05-07 · the 9% is probably higher




FRAMEWORK ARC

How the framework reached production

Four stages over eighteen months. The connective tissue is called "momentum."





FINDING 01 • STRUCTURED INTAKE

Structured intake performed above expectations — volume and latency were not the problems we thought

What worked


API connectors handled 94% of structured signals with no manual touch, at 4-minute average latency. This is the part of the demo we show everyone.

What required tuning

NPS verbatim classification ran an 18% error rate for the first two weeks. It's mentioned in the appendix, not in the headline numbers.

KEY INSIGHT

Viable as designed. The headline numbers are accurate. The denominator is carefully chosen.



Three scoring failure modes found in the pilot

- 1

Recency dominance


Recency set above 50% — last week felt more real than six months of NPS. Capped at 40%, until a VP overrides it.
- 2

Source concentration


One customer was 34% of intake. Also 31% of revenue, so the diversity floor waited for renewal.
- 3


Outcome misclassification

"Improve retention" won't score. "Make customers happier" was submitted twelve times, scored 5 each.





Where the 18 logged decisions landed






The reasons logged for not adopting the framework





A product manager's first month with the framework





WALK THESE QUESTIONS WITH ME IN 60-90 MINUTES. THE OUTPUT IS A DESIGN WE CAN EXECUTE — OR
AGREEMENT THAT WE NEED ANOTHER SESSION TO DESIGN IT.

WHAT WOULD HELP US MOVE FORWARD

**Next step is a working
session, not a debate**