

EXECUTIVE BRIEFING • 2026-08-25

Not Another Presentation App

What is actually here, ranked by what it is worth — and what it is not.



The category error is thinking this is about slides.

Ten seconds of looking versus ten minutes of reading

WHAT IT LOOKS LIKE

Markdown goes in, a slide comes out. Another take on a thirty-year-old idea, with fewer buttons than the incumbent and no drag handles.




WHAT IT IS

A compiler with a typed component contract, a color system that repairs its own contrast, and seventy-two machine gates standing between an author and a broken artifact.

The slides are the output format. They are not the product.

COUNTED FROM THE REPOSITORY AT HEAD, 2026-08-25

What is behind the thing that looks simple.



This briefing has two tiers, graded differently.

Visibility decides the treatment, not importance — the second tier is where the defensibility lives.

01 **The marqu...** Ten things you can see and check in an afternoon.

02 **The substr...** The engineering nobody sees and everything rests on.

READ THIS SLIDE TWICE

This deck is a demonstration of itself.

One source file. Six reader views, and the reader picks. The full deck is 52 slides; whichever view you are holding came out of the same file.

The same file carries a **bottom line**, an **investor**, a **buyer**, a **board** and a **three-page** view — different slides, same document, no copies. Membership is bound to a content hash, so editing a slide silently de-approves every view until a human re-approves it. No other presentation tool ships this.

What this covers.

01 The marquee — ten things you can see

.....

02 The substrate — the engineering underneath

.....

03 What it cannot do, named plainly

.....

04 The ask

.....

TIER ONE

The Marquee

MEASURED AGAINST MARP, SLIDEV, BEAMER AND QUARTO

Four capabilities with no equivalent in the tools it competes with.

01	Reader lenses	One source file, many audience views, membership bound to a content hash	UNIQUE
02	Honest clipping	The only one that reports the loss on the artifact rather than in a log	UNIQUE
03	Achromatopsia support	Five color-vision palettes, with texture carrying meaning where color cannot	UNIQUE
04	Word-level narration	The deck reads itself aloud in sync, diagrams included	UNIQUE

5

One deck carries five audience views, and the reader picks. No duplicate files, no divergent copies.

Membership is a diff

A slide is tagged only where it differs from its view's base, so the source stays clean.

Approval is a content hash

Not a boolean. Edit any slide and every view de-approves itself until re-approved.

It fails closed

An unapproved or drifted view shows nothing. A cooperating renderer will never over-show.

VERIFICATION

The fail-closed property, run against the shipped library.

node -e "lensEligibility(slides, reg, id)" — this deck's own front matter

```
full      {"status":"ok"}
brief     {"status":"unavailable","reason":"unapproved"}
investor  {"status":"unavailable","reason":"unapproved"}
buyer     {"status":"unavailable","reason":"unapproved"}
board     {"status":"unavailable","reason":"unapproved"}
pager     {"status":"unavailable","reason":"unapproved"}
nope      {"status":"unavailable","reason":"unknown"}
```

— Two refusal reasons here, five in the contract, and none of them is "show the full deck instead." The identity view is the only one available by default, and it is safe because it is everything.

WHY THIS IS HARD TO COPY

The interesting part is what it refuses to do.

Anyone can filter slides. The engineering is in refusing to guess.

The suggester proposes membership from a transparent rule table and **writes nothing**. The reader path never imports the suggester. The only bridge between them is a human pressing Approve, which stamps the hash.

The hash is unkeyed, so it detects drift, not forgery: anyone who can write the source can recompute it. The assurance is the Approve gate — a person looked — not a cryptographic property. This deck states it that way because the project retired the stronger wording in a dated correction, and a deck about measuring yourself does not get to quote the claim its own record withdrew.

MARQUEE

Accessibility Is Not A Checkbox Here

PALETTES SHIPPED FOR COLOR VISION DEFICIENCY

5

Protanopia, deuteranopia, tritanopia, achromatopsia, and a neutral base.

Achromatopsia is total color blindness — roughly one in thirty thousand people, and almost universally ignored by design systems. Lattice ships a palette for it and a gate that keeps it honest.

Color is never the only thing carrying meaning.

Texture as a channel

Every categorical slot owns a pattern as well as a hue, so a chart survives being printed in grayscale or read by someone who sees none of it.

Redundant encoding

Status is a glyph and a position and a label, not only a color. The check, the dash and the cross differ in shape before they differ in ink.

Contrast repaired, not checked

The engine derives every pair in perceptual color space and repairs it to the standard, in both light and dark, rather than warning you afterward.

Most tools check the stylesheet. This one measures the pixels.

Why that distinction decides whether the claim is true

CHECKING THE CSS

Reads the declared colors and compares them. It cannot see a translucent overlay, a gradient, a composited card on a tinted band, or a value that a later rule overrode.



MEASURING THE RENDER

Drives a real browser, reads the painted result, and computes contrast from what a human would actually see. Composition defects only exist on this tier.

The second one found defects the first had certified as passing for months.

THE PART THAT MAKES IT STICK

An accessibility property nobody enforces decays in a quarter.

Nine of the seventy-two gates exist to keep this true.

They check that every palette declares its status trio, that categorical ink is declared and has a fallback, that muted tiers clear their floors, that syntax highlighting stays legible, and that a theme cannot register without the tokens the contract requires. A palette that regresses does not ship — the gate fails the build, not a code review.

MARQUEE

The Only One That Admits It Clipped

THE FAILURE EVERY TOOL HAS AND ONLY ONE REPORTS

Overflowing content does not warn you. It just disappears.

Put twelve items on a slide sized for six and every tool renders something. Four of the five we tested drop the surplus and exit successfully. The deck looks finished. The missing risk item is discovered in the room.

Lattice measures the rendered box against its container, and when content will not fit it says so — a visible marker on the slide and a named page on standard error.

SAME BRIEF, SAME TWELVE ITEMS, FIVE TOOLS, MEASURED 2026-08-23

What each tool does when the content does not fit.

TOOL	ITEMS REACHING THE PAGE	HOW YOU FIND OUT
Lattice	5 of 6 retained	Marker on the slide, named page on stderr
Beamer	0 of 6	One warning line, if you read the log
Marp	2 of 6	Nothing
Slidev	2 of 6	Nothing
Quarto	not comparable	Nothing

— Silent truncation is the industry default. Being loud about your own failure is the differentiator.

THE HONEST VERSION OF THIS CLAIM

The mechanism is narrower than the marketing would like.

Automatic re-pagination is deliberately off at widescreen.

At 16:9 an overstuffed deck clips and reports; at other sizes the engine splits the slide and adds pages. That was a design decision about what a presenter expects, not a gap — but it means the headline is "it tells you," not "it fixes it." An earlier draft of this claim said the wrong thing and measurement corrected it.

Word-level synchronization

The highlight tracks the spoken word, not the slide, so a reader can follow along or skim ahead.


Diagrams narrate too

A flowchart is described in reading order rather than skipped as an image.

The timing is estimated, not guessed

Trailing silence and spoken number length are modeled, because "twenty twenty-six" is not two syllables.

MARQUEE · NARRATION

A large, bold, black letter 'O' is centered on the page. Below it is a horizontal line.

Words a presenter must record for a deck to narrate itself, aloud, in sync.

MARQUEE

**The Model Proposes.
The Kernel Disposes.**

THE ARCHITECTURAL BET

A language model is never allowed to author the artifact.

It authors a proposal. Deterministic code turns the proposal into output.

Ask for a theme and the model returns ten colors — not a stylesheet. The engine derives the rest and repairs every pair to standard. Ask for a component and it returns a manifest and CSS that must clear the same gates a first-party component clears, or it does not render. The failure mode of a confident wrong answer becomes a rejected proposal instead of a broken deck.

Four faculties, one discipline.

Themes

Ten colors in, a full audited palette out, repaired in both canvases.

Components

A draft that must clear the token, scope, margin and manifest gates before it renders once.

Finishes

A closed vocabulary of four layers. Out-of-vocabulary values clamp rather than leak.

Motion

The model emits data describing a scene, never code that runs.

WHERE THIS IS GENUINELY UNPROVEN

We test the disposer. We have never tested the proposer.

Every one of those four faculties is unit-tested against a reply we wrote by hand.

That proves the kernel survives a well-formed proposal. It does not prove any real model produces one. Exactly one faculty has a live-model evaluation, against a single pinned model that is now out of date. This gap is filed, scoped and unstarted — and it is the honest boundary of the claim on the previous slide.

61

Contracted components across thirteen buckets. Not a blank canvas with a template gallery bolted on.

A component can declare its capacity

How many items it holds before it crowds, and the count past which it overflows. Nineteen of sixty-one do; the rest is unfinished work, named again on the disclosure slides.

Every component declares its shape

Which form it is, what substance it carries, and which component to escalate to when the count blows the budget.

The catalog is the contract

A component that does not satisfy its manifest cannot register, so the catalog cannot drift from what ships.

SYNTHESIZED FROM REYNOLDS, DUARTE, MINTO AND KNAFLIC; THE THRESHOLDS ARE OURS

The engine has an opinion about how much you may write.

- 01 Slide title 10 words, hard stop at 14
- 02 Subtitle 12 words, hard stop at 18
- 03 Key insight 18 words, hard stop at 28
- 04 Whole slide 70 words and 6 bullets, whichever comes first

WHY A BUDGET AND NOT A WARNING

Verbosity renders perfectly. That is exactly the problem.

A wall of text is not a rendering failure, so no renderer has ever caught it.

These budgets are routed to the review tier rather than the error tier — deliberately. Overrunning them is a communication defect, not a broken build, and the tool says so in that voice. It is the only presentation engine that ships an opinion about prose length at all, and the opinion is cited rather than asserted.

The author picks ten

Canvas, ink, accent, and the status trio. Nothing else.

The engine derives the rest

In perceptual color space, so a lightness step means the same thing at every hue.

Twelve categorical slots

Each a fill and a mark that flip together between canvases, so a chart stays legible in either mode without a second palette.

10 → 107

Ten chosen colors become a hundred-and-seven-token contract, repaired to standard in both light and dark.

Two libraries most decks never need, and some decks cannot live without.

Scenes, not animations

A motion scene is authored as data in a closed vocabulary and validated before it plays. There is no scripting surface to get wrong.

Self-driving walkthroughs

A product tour that drives the real interface, so the demo cannot drift from the thing it demonstrates.

Diagrams that draw themselves

A process flow renders as vector and reveals in reading order rather than appearing all at once.

Charts reveal, then settle

Motion carries the sequence of an argument instead of decorating it.

One source, four destinations, none of them lossy by accident.

Vector PDF.

Real text, selectable and searchable, with fonts embedded so it renders identically on a machine that has none of them.

PPTX.

For the room that will not accept anything else, generated rather than hand-rebuilt.

Standalone HTML.

A single file with no network dependency — it opens on a laptop with the wifi off.

A player.

The deck as a self-contained interactive artifact, pruned to only the styles it actually uses.

THE PROPERTY THAT MATTERS TO A COMPLIANCE READER

The same input produces the same bytes.

Render this deck twice and the two files are identical.

Timestamps are the usual reason a document cannot be diffed, so they are rewritten in place. That makes a rendered deck reviewable the way source is reviewable — you can prove which bytes changed between two versions, and a build that was supposed to change nothing can be shown to have changed nothing.

The authoring
surface is a
browser tab,
and it runs the
real engine.

Not a preview

The same kernel that renders the export renders the editor, so what you approve is what ships.

The model is bring-your-own-key

AI assistance runs on the reader's own credentials. No server-side key, no per-seat inference cost.

It works offline

The engine, the catalog and the themes are a static bundle. The network is for saving, not for rendering.

TIER TWO

The Substrate

NAMED CHECKS STANDING BETWEEN AN AUTHOR AND A BROKEN ARTIFACT

72

One ten-thousand-line file, run on every push. A handful encode house style; the rest exist because something broke, was root-caused, and the cause became a thing the machine now refuses.

THE SECURITY ARCHITECTURE, IN ONE PARAGRAPH

Untrusted markdown reaches a live frame through five different doors.

Closing four of them is the mistake that gets made.

A preview builder sanitizes markup — and then embeds theme stylesheets two lines below, which a browser parses as raw text where a stray closing tag ends the element and everything after it becomes live markup. A separate arm covers that. A third covers modules that take stylesheets back out of a document and re-wrap them, because the serializer re-normalizes the escape. A fourth covers markup injected inside the frame after sanitizing. Each door was found the hard way and each is now a gate.

Determinism, and why it is not a nice-to-have.

✓	Rendered bytes are reproducible — timestamps rewritten, output diffable	shipped
✓	Font metrics pinned, so text does not reflow between machines	shipped
✓	Committed artifacts gated for freshness against their source	shipped
✓	Line endings and byte-order marks normalized at every ingest boundary	shipped
✓	Null bytes refused at the gate rather than corrupting a render	shipped
—	Timing benchmarks split into two signals — workload gates, wall-clock reports	by design
○	Wall-clock timing comparable across machines	not possible

— The last row is honest: the same code has read 93.9ms and 43.1ms on one runner. The benchmark refuses to gate on a number it cannot trust, which is why the row above it exists.

THE MEASUREMENT TIER — WHERE CLAIMS GO TO DIE

Five things the project measures that most projects assert.

01	Rendered contrast	A real browser paints the slide; contrast is computed from the pixels	FINDS COMPOSITION DEFECTS A STYLESHEET AUDIT CANNOT SEE
02	Overflow	Every slide's content box measured against its container	THE BASIS OF THE HONEST-CLIPPING CLAIM
03	Cross-renderer equivalence	The same deck through two paths, compared	CATCHES A TRANSFORM THAT ONLY ONE PATH APPLIES
04	Committed goldens	362 reference PDFs, ~150 MB, diffed on demand	NO CADENCE — AN ACKNOWLEDGED GAP
05	Engine throughput	Head against base on one runner, nightly	THE ONLY TIMING SIGNAL THAT IS GATED

THE PART BUILT FOR MACHINES RATHER THAN PEOPLE

The catalog is designed against a token budget.

An agent choosing a component should not have to load the catalog to choose one.

So the catalog publishes a pick surface: one line per component, the whole set in a few thousand tokens, enough to choose and nothing more. The full machine record is a separate artifact that tools read and authors never load. This is an economics decision about the cost of every future authoring session, and it is the kind of thing that only shows up in a bill.

Engineering process, enforced rather than encouraged.

01	Decision records	463 dated notes, 296 shipped, each one a root cause with an index that refuses to drift	ENFORCED
02	Changelog fragments	One file per change, folded in at release — a shared region that ejected seven merges in one evening	ENFORCED
03	Rebase before push	Folded into the push rather than a background watcher that thrashed the merge train	DISCIPLINE
04	One feature, one branch	Stacked chains banned after they fragmented review	DISCIPLINE
05	Agent model pinning	Every automated reviewer declares its tier; cheaper tiers rejected by name	ENFORCED

THE FIVE LIBRARIES, AND WHY THEY ARE STILL IN-TREE

Five subsystems are already libraries in everything but publication.

Each has a public surface, its own tests, and a gate that stops the engine reaching into it. Four of the five also ship a built distribution.

They cover reader lenses, narration and read-along, audio, motion scenes, and guided walkthroughs. None is published to a registry. That is a deliberate deferral rather than an oversight — but it means the standalone-value argument is currently a design claim, not a download count.

TIER THREE

What It Cannot Do

READ THIS BEFORE THE MARQUEE CONVINCES YOU

Every number in this deck is measured. Several of them are unflattering.

A pitch that names no limits is asking you to trust the parts you cannot check.

What follows is the same audit the engineering team runs on itself, unedited. It is here because the differentiator in everything above is that this project measures itself — and a deck that claimed otherwise would be evidence against its own argument.

The four hardest questions, answered before they are asked.

01 Is any of this actually adopted?

No. No users, no revenue, no external deployment. That is what the ask is for: the engineering is real and you would be the first evidence that it matters.

02 What is the bus factor?

One. A single primary author, and a manual that assumes agent throughput rather than a team.

03 Can we build on it commercially?

Not without a conversation. AGPL-3.0-only is deliberate, and a real constraint on embedding.

04 Are the libraries reusable?

Unproven. Five have public surfaces. None is published, so nobody outside this repo has installed one.

Claims this deck deliberately does not make.

☐	Real-device behavior on phones — marked unverified throughout, never tested on hardware	unverified
☐	"Faster than the alternatives" — the benchmark refuses to compare across machines	unmeasurable
☐	AI output quality — the kernel is tested, the models never were	unstarted
☐	Per-component capacity coverage — under half the catalog declares a budget	partial
☐	Golden image freshness — 362 references, ~150 MB, no cadence watching them	ungated
☒	Everything on the marquee slides, each with an artifact behind it	measured

— Six rows, and only the last one is good news. That ratio is the point of the slide.

Three places the architecture leaks, named plainly.

01	The style escape hatch	Front matter can override any derived token. A deck was reproduced at 1.06-to-1 contrast, exiting successfully with a clean lint	OPEN
02	Export to the original renderer	The compatibility path is broken and has been for some time; the independence story is real, the retreat path is not	BROKEN
03	Documentation drift	A canonical doc contradicted the generated catalog for months, and was corrected only because this audit named it	RECURRING

— The first is a design decision with a missing guard, not a bug. Experienced authors should be able to override. Models should not, and today nothing distinguishes them.

What this is, stated without the adjectives.

The two-sentence version for someone who reads no further

WHAT IS PROVEN

A rendering engine with unusual depth in accessibility, determinism and machine-checkable invariants, verified by seven thousand tests and seventy-two gates that run on every push.



WHAT IS NOT

That anyone wants it. Every claim in this deck is about the artifact. None of them is about a market, a user, or a dollar.

CLOSE

The Ask

THE ASK

Try it. Then decide whether it is worth anything.

STEP 01

Open

lattice.style

— the Studio runs the real engine in a browser tab. No install, no account, no key.



STEP 02

Make one real deck.

Something you actually have to present, not a demo.



STEP 03

Decide. If it is worth something to you, that is the first evidence this project has ever had.

— Everything else in this deck is an argument that the hour is worth spending. Nobody has spent it yet, which is the whole of the ask and the reason the disclosure slides read the way they do.

WHAT TO TAKE AWAY

It is a compiler that happens to emit slides.

The simple surface is the achievement, not the disguise.

Everything on the marquee exists so that an author does not have to think about it, and everything in the substrate exists so the marquee cannot quietly stop being true. What is unproven is whether anyone wants that — which is the only question this deck cannot answer.