

LATTICE • PORTRAIT PROSE DE-BOOST

Dense layouts that fit the tall frame.

The same content-dense layouts, on a 9:16 frame. Body prose de-boosts so the title and last item stay on the slide — hero type keeps its presence.

The framework has four components.

Signal Intake

Weekly structured collection across conversations, market data, and competitive moves.

Scoring Model

Each signal scored on confidence, recency, and relevance. Weights reviewed quarterly.

Decision Log

Every decision recorded with the signals that informed it and the criteria applied.

Calibration Loop

Monthly retrospective comparing predicted to actual outcomes, adjusting the weights.

Who owns each part of the lifecycle.

Signal custody

Signal owner

Runs intake quality. Never tunes the weights — only picks which signals surface.

Policy

Framework operator

Owns scoring policy, calibration cadence, and the rollback playbook. One person.

Consumption

Product team

Holds scoring profiles; runs intake and decision-logging; finds the bugs first.

Oversight

Auditor

Reads the audit trail, cannot edit weights, and is the only role anyone fears.

Four failure modes the framework prevents.

False signal amplification

One loud voice dominates the decision. The model caps any single source at 30% of weight.

Signal hoarding

Teams collect signals but never log decisions, so the calibration loop has nothing to learn from.

Recency bias drift

The newest signal feels most urgent; without recency weighting it overrides calibrated history.

Confidence inflation

Owners self-mark signals high-confidence; the model audits that against realized outcomes.

Signal value versus effort to capture.

High value, low effort

Win/loss interviews. Already happening; just needs structured capture.

High value, high effort

Longitudinal cohort tracking. Worth it, but staff it deliberately.

Low value, low effort

Sentiment tags. Cheap to collect, rarely change a decision.

Low value, high effort

Open-ended surveys. Expensive to analyze, low signal density.

What we recommend, and why.

BUILD

Full control, but 2–3 engineer-quarters and ongoing maintenance.

BUY

Ship in 6 weeks; redirect engineering to product-layer features.

PARTIAL

Hybrid splits ownership and doubles the integration surface.

DEFER

Costs a quarter of signal data we can never reconstruct later.

The scoring model, before and after calibration.

BEFORE CALIBRATION

Equal weights across all three dimensions. Confidence, recency, and relevance each contribute 33%. Simple and consistent, but blind to what your market rewards.



AFTER CALIBRATION

Weights reflect your team's historical accuracy. If recency has been the weakest predictor, it gets downweighted. The model becomes a record of what you have learned.

The shift takes two retrospective cycles to stabilize.

Build the data layer or buy it?

Both paths are viable. The difference is where we spend the next 18 months.

BUILD IN-HOUSE

- Full control over schema and roadmap
- 2–3 engineer-quarters to feature parity
- Ongoing maintenance stays internal

BUY + CONFIGURE



- Ship in 6 weeks, not 9 months
- Engineering redirects to product features
- Exit risk manageable via data export

RECOMMENDATION

Buy the infrastructure. Build the differentiation.

LATTICE • PORTRAIT PROSE DE-BOOST

Hero type
stays; dense
prose adapts.