

QUADRANT • PER-SLIDE LABEL SIZING


One crowded slide set the type for all of them.

A quadrant prints each item's name beside its dot. That name had **one size**, written once in the kernel and shared by every quadrant in every deck — and the number had to be the largest the most crowded slide in the corpus could carry. So a four-dot chart with a half-empty plot printed at the size a fourteen-dot chart needed, in a different file, written by someone else.


The constant was a floor disguised as a size.

`FS.dotLabel` sat at **9.5** because the fourteen-initiative stress slide loses a name at 10.5 — which is exactly what happened once, in #1544, and it reached a committed gallery PDF before anyone caught it. The number was correct for that slide and wrong for the twenty-five others paying for it. Measured across the corpus: **twenty-five of twenty-seven** quadrant slides were running below what their own geometry allows, and not one reached the house's smallest type tier.

Where to put the next dollar.



Ten initiatives, and every name still reads.



Ten items – still clears the bar, because names are short


The size is asked of the geometry, not chosen.

The kernel walks a ladder of sizes **largest first** and keeps the first rung where every name still places. `placeLabels` is pure and deterministic, so this is a measurement, not a heuristic. The bottom rung is the old constant — a slide the ladder cannot lift renders **byte-identically** to before.

The box grows with the type, and that was not optional.

`LW.item` is a width in `viewBox` units. It buys *characters* only in combination with a font size — held at 120 while the type grew, a line fell from about twenty-one characters to fifteen, and names began breaking mid-word: `Maturity self-assessment` rendered as **self-assessmen / t**. Scaling the budget by the same factor holds characters-per-line constant, so a label is a larger copy of itself.

The density that pins the floor.



The ceiling is a design decision, not a limit we hit.

Plenty of sparse quadrants place every name at 18 or 20, and they are **not allowed to**. A dot name that outgrows `--quadrant-axis-size` inverts the chart's hierarchy — the thing being measured would print larger than the axis it is measured against. The ladder stops at the house's smallest type tier: reaching the legibility bar is the goal, and passing it is a different defect.

What this does not fix.

Six slides still cannot carry a readable name, and count is only half of why — name **length** and dot proximity bind too, so some six-item slides settle a rung below every ten-item slide. Those keep the floor size: no worse, no better. Warning an author *before* they build a fourteen-dot quadrant needs a signal the chart family does not have; it is tracked separately.