

FIT LADDER • SPLIT WHAT CAN'T PAGINATE

The read- across carousel

A comparison can't be cut down the
middle — so when it won't fit, it
becomes a sequence.

What the carousel solves

A list overflows? Split it
between items



What the carousel solves

A list overflows? Split it between items — that is **auto-split**

A comparison reads across its two sides →

What the carousel solves

(cont.)

A comparison reads *across* its
two sides — slice it and the
meaning is gone

So the engine re-authors it as a sequence →

What the carousel solves

(cont.)

So the engine **re-authors** it as a sequence: a cover, one page per side, a verdict

Each frame stands alone, reads in order →

What the carousel solves

(cont.)

Each frame stands alone,
reads in order — no setting,
no shrinking

Scoring model: before and after the calibration loop

Side by side →

Before Calibration

Equal weights — confidence, recency, relevance at 33% each. Honest about the guessing. Unchanged in eighteen months.

After Calibration →

After Calibration

Weights reflect each team's accuracy — but the team keeps changing, and the history is one atypical quarter.

**The evaluation came
down to a question
the vendors could not
answer**

Side by side →

The evaluation came down to a question the vendors could not answer

Buy a vendor framework

Three evaluated. None expose calibration weights — the criterion that eliminated all three, added after the team chose to build.

Build the framework in-house →

The evaluation came down to a question the vendors could not answer

Build the framework in-house

Owens the policy, the loop, the timeline. The window closes in eighteen months; a cutover eats nine.

the note →

The evaluation came down to a question the vendors could not answer

The conclusion came first; the evaluation was built to hold it. Everyone in the room knew it.

Decisions used to require a quarterly re-litigation

Side by side →

Before

Decisions lived in the room they were made in. Six months on, nobody could say why a project was killed.

After →

After

Every decision is logged with its signals, options, and bet. Teams still relitigate, but the argument is shorter.

SECTION 01 · DEEP DIVE

Scoring Model Deep Dive

The most configurable component. Six dimensions:

The supporting detail →

Confidence

Corroborating sources, 1–5. Enterprise counts as 1.

Recency →

Recency

Time-decay, configurable half-life. Set to two weeks.

Strategic Relevance →

Strategic Relevance

Owner-scored, 1–5. A 5 tracks whoever presents the roadmap.

Source Diversity →

Source Diversity

Rewards breadth — in practice, newsletter volume.

Volatility →

Volatility

Penalizes swing — including the signal that caught the last surprise.

Owner Bias →

Owner Bias

The correction nobody applies, scored by the owner it is meant to correct.

The six signal dimensions, what they measure, and how they score

Entry by entry →

The six signal dimensions, what they measure, and how they score

Confidence

Independent corroborating sources

1–5, enterprise counts as one

Recency →

The six signal dimensions, what they measure, and how they score

Recency

Time-decay on a configurable half-life

Two-week default surprises everyone

Strategic Relevance →

The six signal dimensions, what they measure, and how they score

Strategic Relevance

Owner-scored against the roadmap

1–5, and a 5 is whoever presents

Source Diversity →

The six signal dimensions, what they measure, and how they score

Source Diversity

Rewards breadth of corroboration

Counts newsletters as sources

Volatility →

The six signal dimensions, what they measure, and how they score

Volatility

Penalizes signals that swing

Also penalizes the early-warning ones

Owner Bias →

The six signal dimensions, what they measure, and how they score

Owner Bias

The correction nobody applies

Scored by the owner it corrects

DECISION · 2026 Q1

**We are
building, not
buying**

The reasoning →

We are building, not buying

Build

Owns the scoring policy, the calibration loop, and the timeline — plus the maintenance, the on-call, and every future explanation of why the framework scored the wrong thing.

Why not buy →

We are building, not buying

Why not buy

Three vendors, none exposing calibration weights. The weights are the product; you cannot buy the product without them. That was the finding.

Why not delay →

We are building, not buying

Why not delay

The window closes in eighteen months — a sentence that has been in this deck, unchanged, since Q1 2025.

BEFORE & AFTER · SCORING MECHANICS

Spreadsheet-driven scoring versus a framework that is basically also a spreadsheet

Both implementations →

```
import pandas as pd

signals = pd.read_csv("signals.csv")
signals["score"] = signals.apply(
    lambda r: 0.33 * r.confidence +
    0.33 * r.recency + 0.33 *
    r.relevance,
    axis=1,
)
signals.to_csv("scored.csv")
```

```
from decision_framework import
Calibrator

cal =
Calibrator.load("weights.json")
signals["score"] =
cal.score(signals)
cal.decisions.log_if_changed(signals
)
```

One finish, more slides

compare-prose / split-panel / list-tabular



One finish, more slides

**compare-prose / split-panel /
list-tabular** open on a cover,
then the content flows on

decision →

One finish, more slides (cont.)

decision makes the verdict the
cover; its reasons flow on

[compare-code →](#)

One finish, more slides (cont.)

compare-code opens on a cover, then one code block per page

Every split wears the same cover finish →

One finish, more slides (cont.)

Every split wears the *same*
cover finish — the deck, just
more of it

More slides, never a clipped comparison

One cover finish across the family — nothing shrinks