FIT LADDER · SPLIT WHAT CAN'T PAGINATE

The read-across carousel

A comparison can't be cut down the middle — so when it won't fit, it becomes a sequence.

What the carousel solves

A list overflows? Split it between items

What the carousel solves

A list overflows? Split it between items — that is **auto-split**

What the carousel solves

(cont.)

A **comparison** reads *across* its two sides — slice it and the meaning is gone

→ next: So the engine re-authors it as a sequence

What the carousel solves

(cont.)

So the engine **re-authors** it as a sequence: a cover, one page per side, a verdict

→ next: Each frame stands alone, reads in order

What the carousel solves

(cont.)

Each frame stands alone, reads in order — no setting, no shrinking

Scoring model: before and after the calibration loop

Side by side →

Before Calibration

Equal weights — confidence, recency, relevance at 33% each. Honest about the guessing. Unchanged in eighteen months.

continues

After Calibration

Weights reflect each team's accuracy — but the team keeps changing, and the history is one atypical quarter.

The evaluation came down to a question the vendors could not answer

Side by side →

THE EVALUATION CAME DOWN TO A QUESTION THE VENDORS COULD NOT ANSWER

Buy a vendor framework

Three evaluated. None expose calibration weights — the criterion that eliminated all three, added after the team chose to build.

continues

THE EVALUATION CAME DOWN TO A QUESTION THE VENDORS COULD NOT ANSWER

Build the framework in-house

Owns the policy, the loop, the timeline. The window closes in eighteen months; a cutover eats nine.

THE VERDICT

The conclusion came first; the evaluation was built to hold it. Everyone in the room knew it.

Decisions used to require a quarterly re-litigation

Side by side →

Before

Decisions lived in the room they were made in. Six months on, nobody could say why a project was killed.

continues

After

Every decision is logged with its signals, options, and bet. Teams still relitigate, but the argument is shorter.

SECTION 01 · DEEP DIVE

Scoring Model Deep Dive

The most configurable component. Six dimensions:

The supporting detail →

Confidence

Corroborating sources, 1–5. Enterprise counts as 1.

next: Recency Time-decay, configurable half-life

Recency

Time-decay, configurable half-life. Set to two weeks.

next: Strategic Relevance Owner-scored, 1–5

Strategic Relevance

Owner-scored, 1–5. A 5 tracks whoever presents the roadmap.

next: Source Diversity Rewards breadth

Source Diversity

Rewards breadth — in practice, newsletter volume.

next: Volatility Penalizes swing

Volatility

Penalizes swing — including the signal that caught the last surprise.

continues

SCORING MODEL DEEP DIVE

Owner Bias

The correction nobody applies, scored by the owner it is meant to correct.

The six signal dimensions, what they measure, and how they score

Entry by entry →

Confidence

Independent corroborating sources 1–5, enterprise counts as one

continues

Recency

Time-decay on a configurable half-life Two-week default surprises everyone

continues

Strategic Relevance

Owner-scored against the roadmap 1–5, and a 5 is whoever presents

continues

Source Diversity

Rewards breadth of corroboration Counts newsletters as sources

continues

Volatility

Penalizes signals that swing Also penalizes the early-warning ones

continues

Owner Bias

The correction nobody applies Scored by the owner it corrects

DECISION · 2026 Q1

We are building, not buying

The reasoning →

Build

Owns the scoring policy, the calibration loop, and the timeline — plus the maintenance, the on-call, and every future explanation of why the framework scored the wrong thing.

continues

Why not buy

Three vendors, none exposing calibration weights. The weights are the product; you cannot buy the product without them. That was the finding.

continues

Why not delay

The window closes in eighteen months — a sentence that has been in this deck, unchanged, since Q1 2025.

BEFORE & AFTER · SCORING MECHANICS

Spreadsheet-driven scoring versus a framework that is basically also a spreadsheet

Both implementations →

BEFORE · THE HONEST SPREADSHEET

```
import pandas as pd

signals = pd.read_csv("signals.csv")
signals["score"] = signals.apply(
    lambda r: 0.33 * r.confidence +
0.33 * r.recency + 0.33 *
r.relevance,
    axis=1,
)
signals.to_csv("scored.csv")
```

```
from decision_framework import
Calibrator

cal =
Calibrator.load("weights.json")
signals["score"] =
cal.score(signals)
cal.decisions.log_if_changed(signals
)
```

One finish, more slides

compare-prose / split-panel / list-tabular

→

One finish, more slides

compare-prose / split-panel / list-tabular open on a cover, then the content flows on

→ next: decision

One finish, more slides (cont.)

decision makes the verdict the cover; its reasons flow on

→ next: compare-code

One finish, more slides (cont.)

compare-code opens on a cover, then one code block per page

→ next: Every split wears the same cover finish

One finish, more slides (cont.)

Every split wears the *same* cover finish — the deck, just more of it (`autosplit: on`)

More slides, never a clipped comparison

One cover finish across the family — nothing shrinks