

FIT LADDER • SPLIT WHAT CAN'T PAGINATE

The read- across carousel

A comparison can't be cut down the
middle — so when it won't fit, it
becomes a sequence.

What the carousel solves

A list overflows? Split it
between items



What the carousel solves

A list overflows? Split it between items — that is **auto-split**

→ next: A comparison reads across its two sides

What the carousel solves

(cont.)

A comparison reads *across* its
two sides — slice it and the
meaning is gone

→ next: So the engine re-
authors it as a sequence

What the carousel solves

(cont.)

So the engine **re-authors** it as a sequence: a cover, one page per side, a verdict

→ next: Each frame stands alone, reads in order

What the carousel solves

(cont.)

Each frame stands alone,
reads in order — no setting,
no shrinking

Scoring model: before and after the calibration loop

Side by side →

Before Calibration

Equal weights — confidence, recency, relevance at 33% each. Honest about the guessing. Unchanged in eighteen months.

→ continues

SCORING MODEL: BEFORE AND AFTER THE CALIBRATION LOOP

After Calibration

Weights reflect each team's accuracy — but the team keeps changing, and the history is one atypical quarter.

**The evaluation came
down to a question
the vendors could not
answer**

Side by side →

THE EVALUATION CAME DOWN TO A QUESTION
THE VENDORS COULD NOT ANSWER

Buy a vendor framework

Three evaluated. None expose calibration weights — the criterion that eliminated all three, added after the team chose to build.

→ continues

THE EVALUATION CAME DOWN TO A QUESTION
THE VENDORS COULD NOT ANSWER

Build the framework in-house

Owens the policy, the loop, the timeline. The window closes in eighteen months; a cutover eats nine.

THE VERDICT

The conclusion came first;
the evaluation was built to
hold it. Everyone in the room
knew it.

Decisions used to require a quarterly re-litigation

Side by side →

DECISIONS USED TO REQUIRE A QUARTERLY RE-LITIGATION

Before

Decisions lived in the room they were made in. Six months on, nobody could say why a project was killed.

→ continues

DECISIONS USED TO REQUIRE A QUARTERLY RELITIGATION

After

Every decision is logged with its signals, options, and bet.

Teams still relitigate, but the argument is shorter.

SECTION 01 · DEEP DIVE

Scoring Model Deep Dive

The most configurable component. Six dimensions:

The supporting detail →

Confidence

Corroborating sources, 1–5. Enterprise counts as 1.

→ next: Recency Time-decay, configurable half-life

SCORING MODEL DEEP DIVE

Recency

Time-decay, configurable half-life. Set to two weeks.

→ next: Strategic Relevance Owner-scored, 1-5

Strategic Relevance

Owner-scored, 1–5. A 5 tracks whoever presents the roadmap.

→ next: Source Diversity Rewards breadth

Source Diversity

Rewards breadth — in practice, newsletter volume.

→ next: Volatility Penalizes swing

Volatility

Penalizes swing — including the signal that caught the last surprise.

→ continues

Owner Bias

The correction nobody applies, scored by the owner it is meant to correct.

The six signal dimensions, what they measure, and how they score

Entry by entry →

THE SIX SIGNAL DIMENSIONS, WHAT THEY MEASURE, AND HOW THEY SCORE

Confidence

Independent corroborating sources 1–5,
enterprise counts as one

→ continues

THE SIX SIGNAL DIMENSIONS, WHAT THEY MEASURE, AND HOW THEY SCORE

Recency

Time-decay on a configurable half-life Two-week default surprises everyone

→ continues

THE SIX SIGNAL DIMENSIONS, WHAT THEY
MEASURE, AND HOW THEY SCORE

Strategic Relevance

Owner-scored against the roadmap 1–5, and
a 5 is whoever presents

→ continues

THE SIX SIGNAL DIMENSIONS, WHAT THEY
MEASURE, AND HOW THEY SCORE

Source Diversity

Rewards breadth of corroboration Counts
newsletters as sources

→ continues

THE SIX SIGNAL DIMENSIONS, WHAT THEY MEASURE, AND HOW THEY SCORE

Volatility

Penalizes signals that swing Also penalizes the early-warning ones

→ continues

THE SIX SIGNAL DIMENSIONS, WHAT THEY MEASURE, AND HOW THEY SCORE

Owner Bias

The correction nobody applies Scored by the owner it corrects

DECISION · 2026 Q1

**We are
building, not
buying**

The reasoning →

WE ARE BUILDING, NOT BUYING

Build

Owns the scoring policy, the calibration loop, and the timeline — plus the maintenance, the on-call, and every future explanation of why the framework scored the wrong thing.

→ continues

WE ARE BUILDING, NOT BUYING

Why not buy

Three vendors, none exposing calibration weights. The weights are the product; you cannot buy the product without them. That was the finding.

→ continues

WE ARE BUILDING, NOT BUYING

Why not delay

The window closes in eighteen months — a sentence that has been in this deck, unchanged, since Q1 2025.

BEFORE & AFTER · SCORING MECHANICS

Spreadsheet-driven scoring versus a framework that is basically also a spreadsheet

Both implementations →

BEFORE · THE HONEST SPREADSHEET

```
import pandas as pd

signals = pd.read_csv("signals.csv")
signals["score"] = signals.apply(
    lambda r: 0.33 * r.confidence +
    0.33 * r.recency + 0.33 *
    r.relevance,
    axis=1,
)
signals.to_csv("scored.csv")
```

AFTER · THE FRAMEWORK

```
from decision_framework import
Calibrator

cal =
Calibrator.load("weights.json")
signals["score"] =
cal.score(signals)
cal.decisions.log_if_changed(signals
)
```

One finish, more slides

compare-prose / split-panel / list-tabular



One finish, more slides

**compare-prose / split-panel /
list-tabular** open on a cover,
then the content flows on

→ next: decision

One finish, more slides (cont.)

decision makes the verdict the
cover; its reasons flow on

→ next: compare-code

One finish, more slides (cont.)

compare-code opens on a cover, then one code block per page

→ next: Every split wears the same cover finish

One finish, more slides (cont.)

Every split wears the *same*
cover finish — the deck, just
more of it

More slides, never a clipped comparison

*One cover finish across the
family — nothing shrinks*