

FIT LADDER • SPLIT WHAT CAN'T PAGINATE

The read-across carousel

A comparison can't be cut down the middle
— so when it won't fit, it becomes a
sequence.

What the carousel solves

- A list overflows? Split it between items — that is **auto-split**
- A **comparison** reads *across* its two sides — slice it and the meaning is gone
- So the engine **re-authors** it as a sequence: a cover, one page per side, a verdict
- Each frame stands alone, reads in order — no setting, no shrinking

Scoring model: before and after the calibration loop

Side by side →

SCORING MODEL: BEFORE AND AFTER THE
CALIBRATION LOOP

Before Calibration

Equal weights —
confidence, recency,
and relevance each
contribute 33%. At least
it is honest that we are
guessing. No team has
changed them in
eighteen months.

After Calibration

Weights reflect each team's historical accuracy — but the team keeps changing, and the history is one atypical quarter. So most teams keep the 33s and call them calibrated.

THE VERDICT

The shift takes two retrospective cycles — 60 days from adoption, or 14 months, depending on who you ask.

**The evaluation came
down to a question the
vendors could not
answer**

Side by side →

THE EVALUATION CAME DOWN TO A QUESTION THE VENDORS COULD NOT ANSWER

Buy a vendor framework

Three vendors evaluated. None expose calibration weights to the customer — the criterion that eliminated all three. It was added to the rubric after the team chose to build.

THE EVALUATION CAME DOWN TO A QUESTION THE VENDORS COULD NOT ANSWER

Build the framework in-house

Owens the scoring policy, the calibration loop, and the timeline. The window closes in eighteen months; a vendor cutover would eat nine of them.

THE VERDICT

The conclusion came first; the
evaluation was built to hold it.
Everyone in the room knew it.

Decisions used to require a quarterly re- litigation

Side by side →

DECISIONS USED TO REQUIRE A QUARTERLY RE-LITIGATION

Before

Decisions lived in the room they were made in. Six months on, nobody could say why a project was killed — only that someone senior had felt strongly.

DECISIONS USED TO REQUIRE A QUARTERLY RELITIGATION

After

Every decision is logged with its signals, its options, and the bet it made. Teams still relitigate, but now there is a record — so the argument is shorter.

SECTION 01 · DEEP DIVE

Scoring Model Deep Dive

The most configurable component — a feature or a warning sign, depending on your team. Six dimensions:

The supporting detail →

Confidence

Independent corroborating sources, 1–5. Enterprise customers count as 1, regardless of volume.

Recency

Time-decay, configurable half-life. Set to two weeks, then surprised only recent news scores.

Strategic Relevance

Owner-scored, 1–5. A 5 correlates with whoever is presenting the roadmap this quarter.

Source Diversity

Rewards breadth. In practice rewards whoever subscribes to the most newsletters.

Volatility

Penalizes signals that swing. Also penalizes the only signal that caught the last surprise.

Owner Bias

The correction nobody applies, scored by the owner it is meant to correct.

The six signal dimensions, what they measure, and how they score

Entry by entry →

THE SIX SIGNAL DIMENSIONS, WHAT THEY MEASURE,
AND HOW THEY SCORE

Confidence

Independent corroborating sources
1–5, enterprise counts as one

Recency

Time-decay on a configurable
half-life Two-week default
surprises everyone

Strategic Relevance

Owner-scored against the
roadmap 1–5, and a 5 is whoever
presents

THE SIX SIGNAL DIMENSIONS, WHAT THEY MEASURE, AND HOW THEY SCORE

Source Diversity

Rewards breadth of corroboration

Counts newsletters as sources

Volatility

Penalizes signals that swing Also

penalizes the early-warning ones

Owner Bias

The correction nobody applies

Scored by the owner it corrects

DECISION · 2026 Q1

**We are
building, not
buying**

The reasoning →

WE ARE BUILDING, NOT BUYING

Build

Owns the scoring policy, the calibration loop, and the timeline — plus the maintenance, the on-call, and every future explanation of why the framework scored the wrong thing.

WE ARE BUILDING, NOT BUYING

Why not buy

Three vendors, none exposing calibration weights. The weights are the product; you cannot buy the product without them. That was the finding.

WE ARE BUILDING, NOT BUYING

Why not delay

The window closes in eighteen months — a sentence that has been in this deck, unchanged, since Q1 2025.

BEFORE & AFTER · SCORING MECHANICS

Spreadsheet-driven scoring versus a framework that is basically also a spreadsheet

Both implementations →

BEFORE · THE HONEST SPREADSHEET

```
import pandas as pd

signals = pd.read_csv("signals.csv")
signals["score"] = signals.apply(
    lambda r: 0.33 * r.confidence + 0.33
    * r.recency + 0.33 * r.relevance,
    axis=1,
)
signals.to_csv("scored.csv")
```

AFTER · THE FRAMEWORK

```
from decision_framework import  
Calibrator  
  
cal = Calibrator.load("weights.json")  
signals["score"] = cal.score(signals)  
cal.decisions.log_if_changed(signals)
```

One finish, more slides

- **compare-prose / split-panel / list-tabular** → a cover, then the content flows on
- **decision** → the verdict is the cover; its reasons flow on
- **compare-code** → a cover, then one code block per page
- Every split wears the *same* cover finish — the deck, just more of it (`autosplit: on`)

More slides, never a clipped comparison

*One cover finish
across the family
— nothing shrinks*