

FEATURE · READ-ALONG CAPTIONS

Every deck can read itself aloud.

Your slides become a timed WebVTT caption track — one `--captions` flag, no audio to record, no key to buy. The words already on the slide ship as an accessible caption file next to the deck.

The narration is already there — on the slide.

A slide already says something. Read-along captions read what is there — the heading, the prose, a chart's computed facts — and lay it on a timeline, word by word, so a viewer can follow along or read instead of listen.

One flag, two kinds of file.

- Deck-level
 - `deck.vtt` — one continuous track across the whole talk, each slide offset by the ones before it.
- Per-slide
 - `deck.01.vtt`, `deck.02.vtt` — each starts at zero, for a player that captions one slide at a time.

The timing is honest, and it is free.

Captions are timed from Cadenza's estimate — a words-per-minute model tuned to how narration is actually paced. No microphone, no text-to-speech key, no cloud round-trip. It runs offline, and it is deterministic: the same slides always yield the same track.

Numbers read the way you say them.

A slide that says `$4.2M`, or `18%`, is timed as four point two million dollars and eighteen percent — spoken as words, not glyphs. The caption shows what you wrote; the timing follows how it is actually said.

Say it your way — the acronym registry.

Add an `acronyms:` block to your front matter and the narration expands the ones you own — `ARR` becomes "annual recurring revenue," `NDR` becomes "net dollar retention" — while the caption still shows the crisp glyph. The author owns the vocabulary, even overriding the built-in dictionary.

This slide's spoken line comes from the front-matter `captions:` map — keyed by this slide's number — which the acronym registry then expands.

Any slide can override what it says.

A `<!-- caption: ... -->` comment is the exact read-as text for one slide. It *replaces* the generated line whole — an override is the caption, not an addition to it.

Speaker notes are not in this chain at all. A note is yours; if a line is meant to be heard, it belongs in a `caption:`.

It rides beside every export.

The `.vtt` is a sidecar, not baked into the deck's bytes. Export a PDF, an HTML player, a PPTX — the caption files sit next to it, ready for any WebVTT-aware player. Your existing artifact is untouched.

One line at the command line.

lattice deck.md deck.pdf --captions

That is the whole feature. Slides in, caption track out — accessible by default, and honest about its timing.