

FEATURE · READ-ALONG CAPTIONS

Every deck can read itself aloud.

Speaker notes become timed WebVTT captions — one `--captions` flag, no audio to record, no key to buy. The narration you already wrote ships as an accessible caption track next to the deck.

The narration is already there — in your speaker notes.

Every slide's `<!-- note -->` is what you'd say out loud. Read-along captions take that text and lay it on a timeline, word by word, so a viewer can follow along or read instead of listen.

One flag, two kinds of file.

Deck-level

`deck.vtt` — one continuous track across the whole talk, each slide offset by the ones before it.

Per-slide

`deck.01.vtt`, `deck.02.vtt` — each starts at zero, for a player that captions one slide at a time.

The timing is honest, and it is free.

Captions are timed from Cadenza's estimate — a words-per-minute model tuned to how narration is actually paced. No microphone, no text-to-speech key, no cloud round-trip. It runs offline, and it is deterministic: the same notes always yield the same track.

Numbers read the way you say them.

A note that mentions four point two million dollars, or eighteen percent, is spoken as words — not glyphs. The caption shows what you wrote; the timing follows how it is actually said.

It rides beside every export.

The `.vtt` is a sidecar, not baked into the deck's bytes. Export a PDF, an HTML player, a PPTX — the caption files sit next to it, ready for any WebVTT-aware player. Your existing artifact is untouched.

LATTICE DECK.MD DECK.PDF --CAPTIONS

One line at the command line.

That is the whole feature. Speaker notes in, caption track out — accessible by default, and honest about its timing.