

AUTO-SPLIT • STRUCTURE DECIDES THE CUT

One thing per slide

A slide holding six things becomes six slides. The markup says how many, so every machine cuts the same deck the same way.

What changed

What changed

- The trigger is **structure**, not a measured overflow — no render decides the page count

→ continues

What changed (cont.)

- A page carries **one** structural element; nothing packs to an authoring budget

→ continues

What changed (cont.)

- Every run opens on a **cover** and closes on the section's **note and key insight**

→ continues

What changed (cont.)

- Each page **names the next one**, so an atomized run still reads as one thing

→ next: the key insight

What changed

KEY INSIGHT

A deck is authored once. What it becomes should not depend on which machine renders it.

What a split page owes the reader

The one thing it holds, set
at full size...



What a split page owes the reader

01 The one thing it holds, set at full size —
not shrunk to share

→ next: A heading that says
which run it belongs...

What a split page owes the reader (cont.)

01 A heading that says which run it belongs to

→ next: A pointer to what comes next

What a split page owes the reader (cont.)

01

A pointer to what comes next

→ next: A way back to the whole – the k-of-N rail...

What a split page owes the reader (cont.)

01 A way back to the whole — the k-of-N rail in the footer band

→ next: An ending that says what the run meant

What a split page owes the reader (cont.)

01 An ending that says what the run meant

→ next: the note

What a split page owes the reader

★ *list-criteria could not split at all before this change.*

How the cut was decided, over time

Opt-in per deck

An author who never heard of the flag got a clipped slide

Default-on for portrait

Right direction, wrong altitude — the directive was the thing to remove

Measured fit

The page count became a property of the renderer


Structure

Knowable from the markup, so the linter and the export agree

Reading a split run

Arrive →

Reading a split run



• reader

PAIN

A chart splits when it has a seam. This one does – independent stages. A pie chart does not.

→ next: Arrive

Reading a split run (cont.)



reader

1

A chart splits when it has a seam. This one does – independent stages. A pie chart does not.

→ next: Arrive

Reading a split run (cont.)



reader

2

A chart splits when it has a seam. This one does – independent stages. A pie chart does not.

→ next: Arrive

Reading a split run (cont.)



reader

3

A chart splits when it has a seam. This one does – independent stages. A pie chart does not.

→ next: Arrive

Reading a split run (cont.)



reader

4

A chart splits when it has a seam. This one does – independent stages. A pie chart does not.

→ next: Arrive

Reading a split run (cont.)



reader

5

A chart splits when it has a seam. This one does – independent stages. A pie chart does not.

→ next: Arrive

Reading a split run (cont.)

Arrive

R

Meet the cover



R

Learn what the run is about



Move

R

Take one element per page



R

Follow the pointer to the next



Close

R

Read the note and the insight together



R reader

DELIGHT

A chart splits when it has a seam. This one does
– independent stages. A pie chart does not.

Enrollment across the catalog

32 of 61 →

Enrollment across the catalog

- **32 of 61** components split — four of them could not before

→ next: 29

Enrollment across the catalog (cont.)

- **29** are single structural elements already, and ring on overflow

→ next: 0

Enrollment across the catalog (cont.)

- **0** pack, at any count, on any component

→ next: Every

Enrollment across the catalog (cont.)

- **Every** run carries a forward pointer; four did before

→ next: the note

Enrollment across the catalog

Each of the 29 has its reason recorded in
`lib/core/split-facts.js`.

More slides, one thing each

Structure decides, not the renderer